

## Assignment 5: Multiprocessor Map-Reduce

### 1. Names of group members.

Hua Yang, 128-00-2637

Erik Kamp, 132-00-4838

### 2. Detailed instructions for using the program, including any preprocessing and postprocessing steps that must be taken. Make sure that you submit only the program source, headers, and makefile. A user should do nothing more than run *make* to compile the program. Make sure your instructions clearly specify the command name and usage and provide an example.

The program will only accept the following command:

```
./mapred -a [wordcount, sort] -i threads -m num_maps -r num_reduces infile output
```

### 3. Architectural description of the framework, including where you had to employ synchronization primitives.

- For this project, we only wrote the multi-threaded version, using the `p_thread` library.
- For splitting the files equal to `x` number of mappers, we use the provided `split.sh` file.
- The structure that we used to store the mappers and reducers output is an `ut_hash`.

- Depending on the specified m number of mappers and r number of reducers, we create m threads of mappers and r threads of reducers.

#### **4. Tests that were run.**

The tests we did utilized the files from Project Gutenberg.

#### **5. Description of difficulties and problems that you encountered.**

- We initially had a problem with forks, where we did not wait for the child processes to finish.
- The threads were tricky because of the race conditions.
- We had trouble coming up with a data structure to store the mapped outputs, turned out that `ut_hash` was the best solution.

#### **6. If you implemented both threads and processes, an evaluation of the performance difference between the two.**

We only did threads, since we're in a group less than 3.